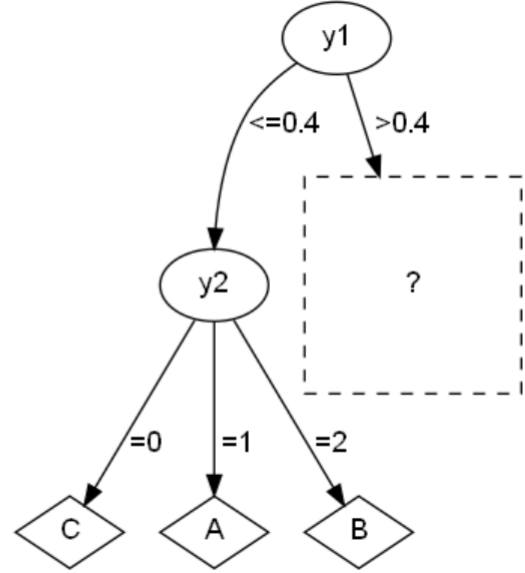


Part I: Pen and paper

Consider the partially learnt decision tree from the dataset D . D is described by four input variables – one numeric with values in $[0,1]$ and 3 categorical – and a target variable with three classes.

D	y_1	y_2	y_3	y_4	y_{out}
x_1	0.24	1	1	0	A
x_2	0.06	2	0	0	B
x_3	0.04	0	0	0	B
x_4	0.36	0	2	1	C
x_5	0.32	0	0	2	C
x_6	0.68	2	2	1	A
x_7	0.9	0	1	2	A
x_8	0.76	2	2	0	A
x_9	0.46	1	1	1	B
x_{10}	0.62	0	0	1	B
x_{11}	0.44	1	2	2	C
x_{12}	0.52	0	2	0	C



1. [5v] Complete the given decision tree using Information gain with Shannon entropy (\log_2). Consider that: i) a minimum of 4 observations is required to split an internal node, and ii) decisions by ascending alphabetic order should be placed in case of ties.

The information gain of variable y_2 , $y_1 > 0.4$ is given by

$$IG(y_2, y_1 > 0.4) = E(y_{out}|y_1 > 0.4) - E(y_{out}|y_2, y_1 > 0.4) \quad (1)$$

The entropy of $y_{out}|y_1 > 0.4$ is given by:

$$E(y_{out}|y_1 > 0.4) = - (p(A|y_1 > 0.4) \log_2 (p(A|y_1 > 0.4)) + p(B|y_1 > 0.4) \log_2 (p(B|y_1 > 0.4)) + p(C|y_1 > 0.4) \log_2 (p(C|y_1 > 0.4)))$$

Since there are **3 As**, **2 Bs** and **2 Cs** we have the following proportions:

$$p(A|y_1 > 0.4) = \frac{3}{7} \quad , \quad p(B|y_1 > 0.4) = \frac{2}{7}, p(C|y_1 > 0.4) = \frac{2}{7}$$

Therefore, we can calculate $E(y_{out}|y_1 > 0.4)$:

$$\begin{aligned} E(y_{out}|y_1 > 0.4) &= - \left(\frac{3}{7} \log_2 \left(\frac{3}{7} \right) + \frac{2}{7} \log_2 \left(\frac{2}{7} \right) + \frac{2}{7} \log_2 \left(\frac{2}{7} \right) \right) \\ &= 1.556 \end{aligned}$$

The next step is calculating $E(y_{out}|y_2, y_1 > 0.4)$:

$$E(y_{out}|y_2, y_1 > 0.4) = E(y_{out}|y_2 = 0, y_1 > 0.4) + E(y_{out}|y_2 = 1, y_1 > 0.4) + E(y_{out}|y_2 = 2, y_1 > 0.4)$$

Then, we can calculate $E(y_{out}|y_2 = 0, y_1 > 0.4)$, $E(y_{out}|y_2 = 1, y_1 > 0.4)$ and $E(y_{out}|y_2 = 2, y_1 > 0.4)$:

$$E(y_{out}|y_2 = 0, y_1 > 0.4) = -\frac{3}{7} \left(\frac{1}{3} \log_2 \left(\frac{1}{3} \right) + \frac{1}{3} \log_2 \left(\frac{1}{3} \right) + \frac{1}{3} \log_2 \left(\frac{1}{3} \right) \right) = 0.67927$$

$$E(y_{out}|y_2 = 1, y_1 > 0.4) = -\frac{2}{7} \left(0 + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) \right) = \frac{2}{7}$$

$$E(y_{out}|y_2 = 2, y_1 > 0.4) = -\frac{2}{7} (1 \log_2 (1) + 0 + 0) = 0$$

Therefore,

$$E(y_{out}|y_2, y_1 > 0.4) = 0.67927 + \frac{2}{7} + 0 = 0.965$$

Calculation of information gain, as per (1):

$$IG(y_2, y_1 > 0.4) = 1.556 - 0.965 = \mathbf{0.591}$$

Then, we do the same for y_3 :

$$E(y_{out}|y_3, y_1 > 0.4) = E(y_{out}|y_3 = 0, y_1 > 0.4) + E(y_{out}|y_3 = 1, y_1 > 0.4) + E(y_{out}|y_3 = 2, y_1 > 0.4)$$

$$E(y_{out}|y_3 = 0, y_1 > 0.4) = -\frac{1}{7} (0 + 1 \log_2 (1) + 0) = 0$$

$$E(y_{out}|y_3 = 1, y_1 > 0.4) = -\frac{2}{7} \left(\frac{1}{2} \log_2 \left(\frac{1}{2} \right) + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) + 0 \right) = \frac{2}{7}$$

$$E(y_{out}|y_3 = 2, y_1 > 0.4) = -\frac{4}{7} \left(\frac{1}{2} \log_2 \left(\frac{1}{2} \right) + 0 + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) \right) = \frac{4}{7}$$

Calculation of information gain, as per (1):

$$IG(y_3, y_1 > 0.4) = 1.556 - \frac{6}{7} = \mathbf{0.699}$$

Then, we do the same for y_4 :

$$E(y_{out}|y_4, y_1 > 0.4) = E(y_{out}|y_4 = 0, y_1 > 0.4) + E(y_{out}|y_4 = 1, y_1 > 0.4) + E(y_{out}|y_4 = 2, y_1 > 0.4)$$

$$E(y_{out}|y_4 = 0, y_1 > 0.4) = -\frac{2}{7} \left(\frac{1}{2} \log_2 \left(\frac{1}{2} \right) + 0 + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) \right) = \frac{2}{7}$$

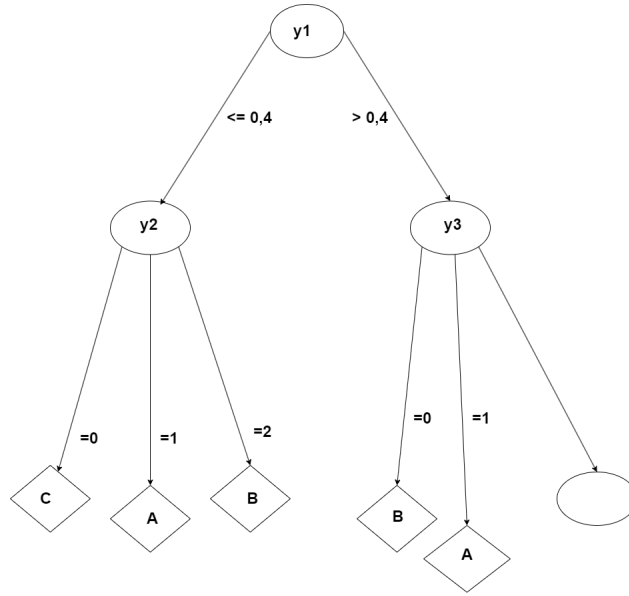
$$E(y_{out}|y_4 = 1, y_1 > 0.4) = -\frac{3}{7} \left(\frac{1}{3} \log_2 \left(\frac{1}{3} \right) + \frac{2}{3} \log_2 \left(\frac{2}{3} \right) + 0 \right) = 0.39356$$

$$E(y_{out}|y_4 = 2, y_1 > 0.4) = -\frac{2}{7} \left(\frac{1}{2} \log_2 \left(\frac{1}{2} \right) + 0 + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) \right) = \frac{2}{7}$$

Calculation of information gain, as per (1):

$$IG(y_{out}|y_4, y_1 > 0.4) = 1.556 - 0.965 = \mathbf{0.591}$$

After the calculations, we know that the Decision Tree will be like this:



To continue the Tree, we need to calculate $IG(y_2|y_3 = 2, y_1 > 0.4)$ and $IG(y_4|y_3 = 2, y_1 > 0.4)$:

$$\begin{aligned} E(y_{out}|y_3 = 2, y_1 > 0.4) = & - (p(A|y_3 = 2, y_1 > 0.4) \log_2(p(A|y_3 = 2, y_1 > 0.4)) \\ & + p(B|y_3 = 2, y_1 > 0.4) \log_2(p(B|y_3 = 2, y_1 > 0.4)) \\ & + p(C|y_3 = 2, y_1 > 0.4) \log_2(p(C|y_3 = 2, y_1 > 0.4))) \end{aligned}$$

$$E(y_{out}|y_3 = 2, y_1 > 0.4) = - \left(\frac{1}{2} \log_2 \left(\frac{1}{2} \right) + 0 + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) \right) = 1$$

In the case of y2:

$$E(y_{out}|y_2 = 0, y_3 = 2, y_1 > 0.4) = -\frac{1}{4} (0 + 0 + 1 \log_2 (1)) = 0$$

$$E(y_{out}|y_2 = 1, y_3 = 2, y_1 > 0.4) = -\frac{1}{4} (0 + 0 + 1 \log_2 (1)) = 0$$

$$E(y_{out}|y_2 = 2, y_3 = 2, y_1 > 0.4) = -\frac{2}{4} (1 \log_2 (1) + 0 + 0) = 0$$

$$IG(y_2, y_3 = 2, y_1 > 0.4) = 1 - 0 = 1$$

In the case of y_4 :

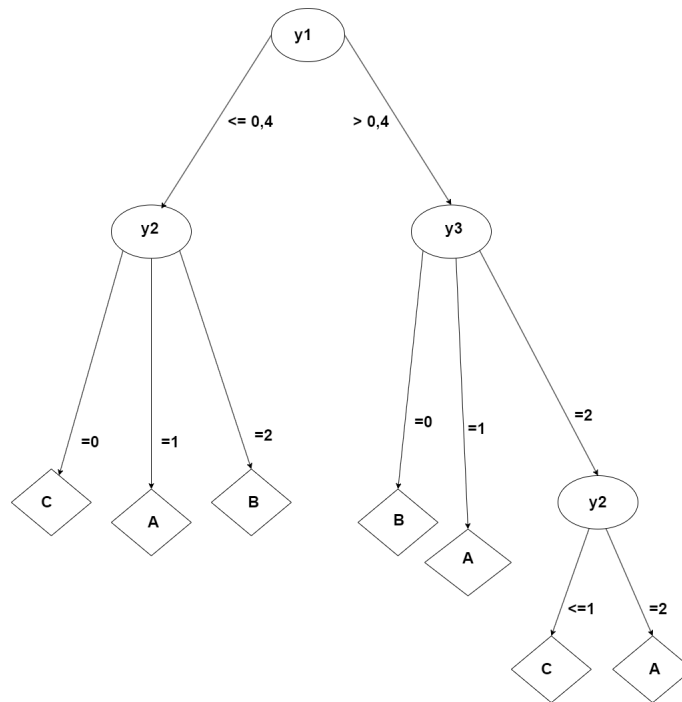
$$E(y_{out}|y_4 = 0, y_3 = 2, y_1 > 0.4) = -\frac{2}{4} \left(\log_2 \left(\frac{1}{2} \right) + 0 + \frac{1}{2} \log_2 \left(\frac{1}{2} \right) \right) = \frac{1}{2}$$

$$E(y_{out}|y_4 = 1, y_3 = 2, y_1 > 0.4) = -\frac{1}{4} (\log_2 (1) + 0 + 0) = 0$$

$$E(y_{out}|y_4 = 2, y_3 = 2, y_1 > 0.4) = -\frac{1}{4} (0 + \log_2 (1) + 0) = 0$$

$$IG(y_{out}|y_4, y_1 > 0.4) = 1 - 0.5 = 0.5$$

Finally, we conclude that the Decision Tree is:



2. [2.5v] Draw the training confusion matrix for the learnt decision tree.

	True A	True B	True C	Total
Pred A	4	1	0	5
Pred B	0	2	0	2
Pred C	0	1	4	5
Total	4	4	4	12

3. [1.5v] Identify which class has the lowest training F1 score.

The training F1 can be calculated by:

$$F_1 = \frac{1}{\frac{1}{2} \times \left(\frac{1}{P} + \frac{1}{R} \right)} \quad (2)$$

Where P is precision and R is recall, which can be calculated by:

$$P = \frac{\text{true positives}}{\text{true positives} + \text{false positives}} \quad R = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}} \quad (3)$$

For class **A**, **4 true positives**, **0 false negatives** and **1 false positive**.

As per (2) and (3),

$$\begin{aligned} F_1 &= \frac{1}{\frac{1}{2} \times \left(\frac{4+1}{4} + \frac{4+0}{4} \right)} \\ &= 0.8889 \end{aligned}$$

For class **B**, we have **2 true positives**, **0 false negatives** and **2 false positives**.

As per (2) and (3),

$$\begin{aligned} F_1 &= \frac{1}{\frac{1}{2} \times \left(\frac{2+2}{2} + \frac{2+0}{2} \right)} \\ &= \frac{2}{3} = 0.6667 \end{aligned}$$

For class **C**, we have **4 true positives**, **0 false negatives** and **1 false positives**.

As per (2) and (3),

$$\begin{aligned} F_1 &= \frac{1}{\frac{1}{2} \times \left(\frac{4+1}{4} + \frac{4+0}{4} \right)} \\ &= 0.8889 \end{aligned}$$

In conclusion, **the class with the lowest training F1 score is B**.

4. [1v] Considering y_2 to be ordinal, assess if y_1 and y_2 are correlated using the Spearman coefficient.

We start by collecting data for variables y_1 and y_2 :

$$y_1 = [0.24, 0.06, 0.04, 0.36, 0.32, 0.68, 0.9, 0.76, 0.46, 0.62, 0.44, 0.52]$$

$$y_2 = [1, 2, 0, 0, 0, 2, 0, 2, 1, 0, 1, 0]$$

Then we do the Ranking of Observations:

$$\text{rank}_{y_1} = [3, 2, 1, 5, 4, 10, 12, 11, 7, 9, 6, 8]$$

$$\text{rank}_{y_2} = [8, 11, 3.5, 3.5, 3.5, 11, 3.5, 11, 8, 3.5, 8, 3.5]$$

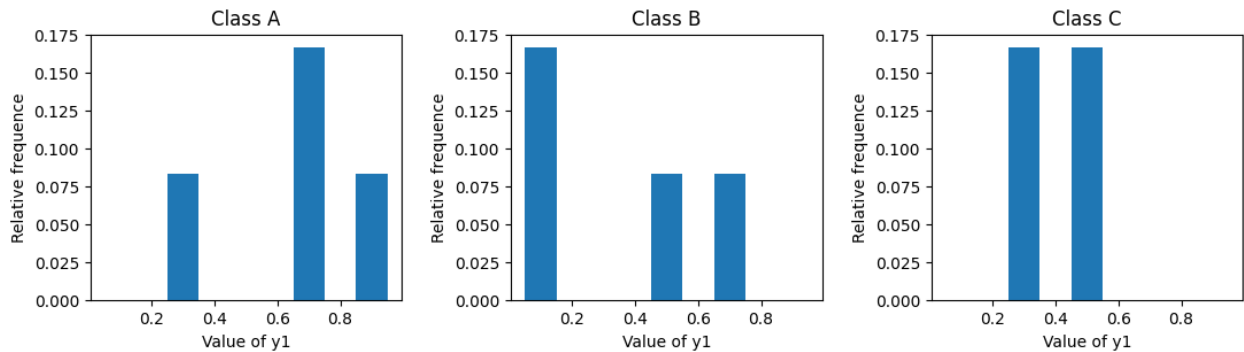
Now, we do the calculation of Spearman's Rank Correlation Coefficient by doing the PCC of the ranks:

$$PCC(\text{rank}_{y_1}, \text{rank}_{y_2}) = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}}$$

$$\begin{aligned} \text{Spearman}(y_1, y_2) &= PCC([3, 2, 1, 5, 4, 10, 12, 11, 7, 9, 6, 8], [8, 11, 3.5, 3.5, 3.5, 11, 3.5, 11, 8, 3.5, 8, 3.5]) \\ &= 0.080 \end{aligned}$$

The calculated value of $Spearman(y_1, y_2)$ provides a measure of the rank correlation (order) between variables y_1 and y_2 . Since the value is **0.080** we can conclude that the variables have a **very weak correlation** between them.

5. [1v] Draw the class-conditional relative histograms of y_1 using 5 equally spaced bins in $[0,1]$. Challenge: find the root split using the discriminant rules from these empirical distributions.



An appropriate root split, given this histogram, would be to slit y_1 into 3 branches:

- $y_1 \leq 0.2$ pointing to B;
- $0.2 < y_1 \leq 0.6$ pointing to C;
- $y_1 > 0.6$ pointing to A.

Part II: Programming

1. Apply `f_classif` from `sklearn` to assess the discriminative power of the input variables. Identify the input variable with the highest and lowest discriminative power. Plot the class-conditional probability density functions of these two input variables.

We used the following program:

```

1 from sklearn.feature_selection import f_classif
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4
5
6 X = df.drop('class', axis=1)
7 y = df['class']
8
9 fimportance = f_classif(X, y)
10 highest_score = fimportance[0].argmax()
11 highest_feature = X.columns.values[highest_score]
12 lowest_score = fimportance[0].argmin()
13 lowest_feature = X.columns.values[lowest_score]
14
15 print('The variable with the highest F-score is', highest_feature, \
16       'with a score of', fimportance[0][highest_score])
17 print('The variable with the lowest F-score is', lowest_feature, \
18       'with a score of', fimportance[0][lowest_score])
19

```

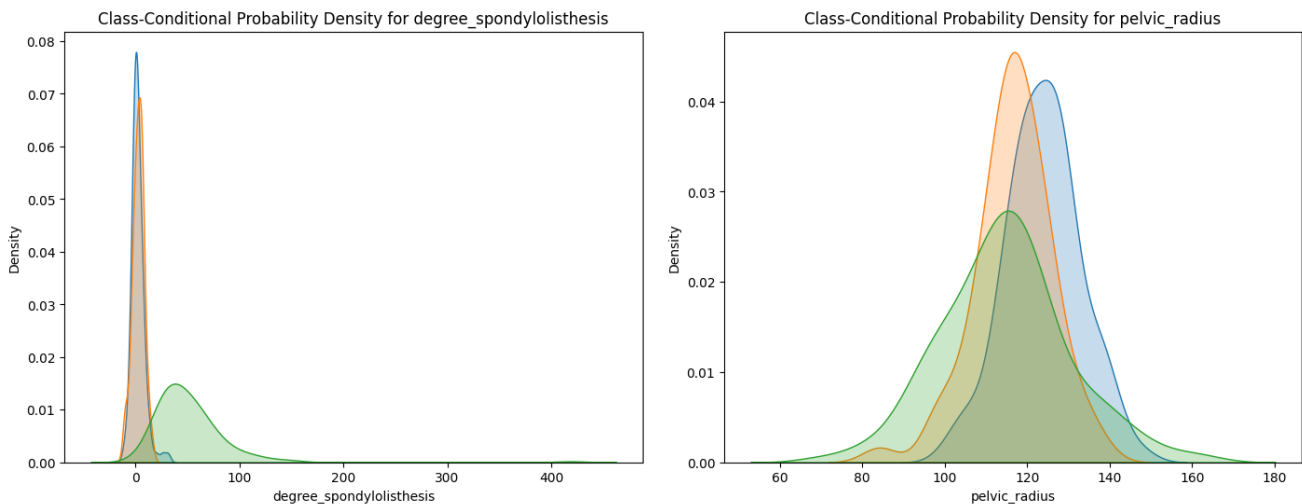
```

20
21 top_2_features = [highest_feature, lowest_feature]
22
23 plot_df = df[top_2_features + ['class']]
24
25 print(plot_df)
26
27 class_hernia_data = plot_df[plot_df['class'] == 'Hernia']
28 print(class_hernia_data)
29 class_normal_data = plot_df[plot_df['class'] == 'Normal']
30 print(class_normal_data)
31 class_Spondylolisthesis_data = plot_df[plot_df['class'] == 'Spondylolisthesis']
32
33 for feature in top_2_features:
34     plt.figure(figsize=(8, 6))
35     sns.kdeplot(class_normal_data[feature], label='Normal', fill=True)
36     sns.kdeplot(class_hernia_data[feature], label='Hernia', fill=True)
37     sns.kdeplot(class_Spondylolisthesis_data[feature], label='Spondylolisthesis',
38                 fill=True)
39     plt.xlabel(feature)
40     plt.ylabel('Density')
41     plt.title(f'Class-Conditional Probability Density for {feature}')
42     plt.show()

```

Listing 1: Plotting the class-conditional probability

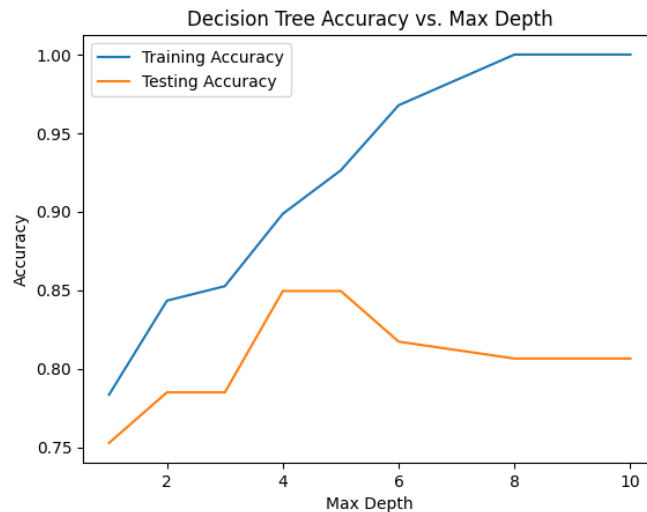
After applying the `f_classif` function we assess that the variable with the **highest discriminative power** is **degree_spondylolisthesis** with a score of **119.12288060759764** and the variable with the lowest is **pelvic_radius** with a score of **16.86693475538006**.



2. Using a stratified 70-30 training-testing split with a fixed seed (random_state=0), assess in a single plot both the training and testing accuracies of a decision tree with depth limits in 1,2,3,4,5,6,8,10 and the remaining parameters as default. [optional] Note that split thresholding of numeric variables in decision trees is non-deterministic in sklearn, hence you may opt to average the results using 10 runs per parameterization.

```
1
2
3 from scipy.io import arff
4 from sklearn.model_selection import train_test_split
5 from sklearn.tree import DecisionTreeClassifier
6 import matplotlib.pyplot as plt
7 from sklearn.preprocessing import LabelEncoder
8
9 data = arff.loadarff('column_diagnosis.arff')
10 df = pd.DataFrame(data[0])
11
12
13 le = LabelEncoder()
14 df['class'] = le.fit_transform(df['class'])
15 X = df.iloc[:, :-1].values
16 y = df['class']
17
18 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
19 random_state=0, stratify=y)
20
21 depth_limits = [1, 2, 3, 4, 5, 6, 8, 10]
22 train_accuracies = []
23 test_accuracies = []
24
25 for depth_limit in depth_limits:
26     train_acc = 0
27     test_acc = 0
28     num_runs = 10
29
30     for _ in range(num_runs):
31         clf = DecisionTreeClassifier(max_depth=depth_limit, random_state=0)
32         clf.fit(X_train, y_train)
33
34         train_acc += clf.score(X_train, y_train)
35         test_acc += clf.score(X_test, y_test)
36
37     train_accuracies.append(train_acc / num_runs)
38     test_accuracies.append(test_acc / num_runs)
39
40 plt.plot(depth_limits, train_accuracies, label='Training Accuracy')
41 plt.plot(depth_limits, test_accuracies, label='Testing Accuracy')
42 plt.xlabel('Max Depth')
43 plt.ylabel('Accuracy')
44 plt.title('Decision Tree Accuracy vs. Max Depth')
45 plt.legend()
46 plt.show()
```

Listing 2: Plotting the class-conditional probability



3. Comment on the results, including the generalization capacity across settings.

Overall, the model seems to perform very well on the training data but poorly on testing data. In the beginning the model seems to perform well since the curves are both similar but when the Depth is near 5 we notice that the Training Accuracy curve increases as the Testing Accuracy decreases showing a significant gap between the training and test performance. This indicates a clear case of overfitting. Therefore we can conclude that the the model has learned to perform well on the training data but struggles to generalize to unseen or new data in different settings so it has a poor generalization and can be considered an unreliable model for a real-world scenario. To improve the model we could review the training data and check if that is any noisy data that does not adequately represent the problem for example. Another possible improvement could be to try to obtain more training data.

4. To deploy the predictor, a healthcare team opted to learn a single decision tree (random_state=0) using all available data as training data, and further ensuring that each leaf has a minimum of 20 individuals in order to avoid overfitting risks. i. Plot the decision tree.

```

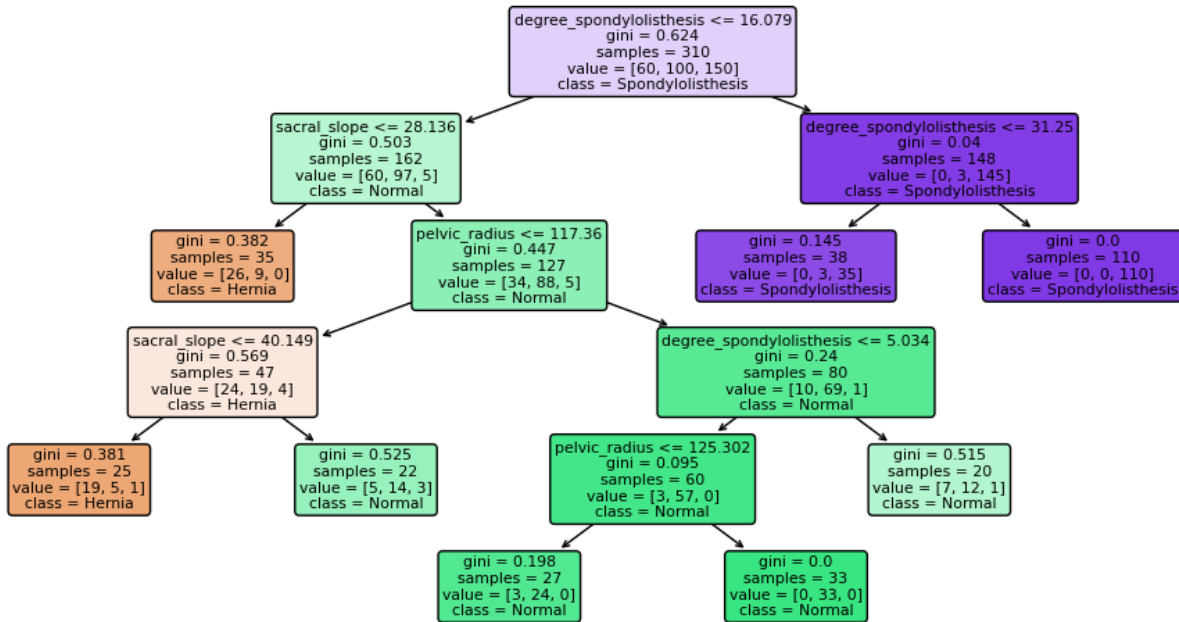
1 import matplotlib.pyplot as plt
2 from sklearn import tree
3 from sklearn.model_selection import train_test_split
4
5 # Load and partition data from column_diagnosis.arff
6 data = arff.loadarff('column_diagnosis.arff')
7 df = pd.DataFrame(data[0])
8 X = df.iloc[:, :-1].values
9
10 # Convert byte strings to regular strings
11 y_str = df['class'].astype(str)
12 # Initialize label encoder
13 le = LabelEncoder()
14 # Fit label encoder and transform labels
15 y = le.fit_transform(y_str)
16
17 # Learn classifier (Decision Tree) with max depth 3
18 predictor = tree.DecisionTreeClassifier(min_samples_leaf=20, random_state=0)
19 predictor.fit(X, y)
20
21 # Plot classifier

```

```

22 plt.figure(figsize=(12, 6))
23 tree.plot_tree(predictor, feature_names=df.columns[:-1], class_names=le.classes_,
24               filled=True, rounded=True)
25 plt.show()

```



ii. Characterize a hernia condition by identifying the hernia-conditional associations

The decision tree shows that having a *degree_spondylolisthesis* ≤ 16.079 and a *sacral_slope* ≤ 28.136 or having a *degree_spondylolisthesis* ≤ 16.079 , $28.136 < \text{sacral_slope} \leq 40.149$ and a *pelvic_radius* ≤ 117.36 probably means you have a Hernia.

End note: do not forget to also submit your Jupyter notebook